

INFS1200/7900

Introduction to Information Systems

Structured Query Language (SQL)

Hassan Khosravi

Notes

- **Lecture:** We will be continuing with Module 4
- **Tutorial:** You will be working through Tutorial 6 on SQL queries. Sample solution for Tutorial 5 has been updated.
- **Practical:** The Practical will focus on providing you with help and guidance on your assignment 1.
- **Contact:** No Contact session will be held during this week.
- **Assignment 1** is due on 21st of September. Make sure that you are using the provided template for answering the questions. We reserve the right to make grade deductions for those that fail to correctly use the provided template. Assignment has room for more creativity compared to exams. State your assumptions.
- **Quiz 2:** Quiz 2 will be held on October 9th. A mock exam for quiz 2 will be released soon.
- **RiPPLE:** Round 3 ends 28 September at 11:59 pm. We have made a few updates to the platform.

Thank you for your feedback on Module 2 and 3

- As per your feedback:
 - I will try to slow down the pace of the lectures.
 - Please feel free to use Piazza if you have any questions about the course content. Thanks to a few brilliant students and their contributions the course has a 64 min average response time.
 - Please attend the lectures, Tutorials and Practicals if you are finding the course challenging.

Learning Objectives

Description	Tag
Create basic SQL queries using: SELECT, FROM, WHERE statements.	
Create SQL queries containing the DISTINCT statement.	
Create SQL queries using set operators.	
Create SQL queries using aggregate operators.	SQL-basic
Create SQL queries containing GROUP BY statements.	
Create SQL queries containing HAVING statements.	
Given a SQL query and table schemas and instances, compute the query result.	
Create nested SQL queries.	
Create SQL Queries that use the division operator.	
Explain the purpose of NULL values, and justify their use. Also describe the difficulties added by having NULLs .	
Create SQL queries that use joins.	
Create SQL queries that use the CHECK statement.	
Create SQL queries that use ASSERTIONS .	
Create, use and modify VIEWS in SQL.	SQL-advance
Modify data stored in a database using the INSERT, DELETE, and UPDATE statements.	
Identify the pros and cons of using general table constraints (e.g., ASSERTION, CHECK) and triggers in databases.	
Show that there are alternative ways of coding SQL queries to yield the same result.	
Determine whether or not two SQL queries are equivalent.	
Create SQL queries for creating tables.	
Create SQL queries for altering tables.	
Create SQL queries to enforce referential integrities.	
Create SQL queries for dropping tables.	SQL-DDL

SQL

- Standardize language, supported by all of the major commercial databases.
- Interactive use via graphical user interface or embedded in programs.
- Declarative, based on relational algebra

Following the Examples

MySQL: an open-source relational database management system

MySQL Workbench: an open source visual database design tool that integrates SQL development, administration, database design, creation and maintenance into a single integrated development environment for MySQL.

Databases used for providing examples: Available for download from the course website.

Movie (MovieID, Title, Year)

StarsIn (MovieID, StarID, Role)

MovieStar (StarID, Name, Gender)

College(cName, state, enrollment)

Student(sID, sName, GPA, sizeHS)

Apply(sID, cName, major, decision)

Borrowed from Rachel Pottinger from UBC

Borrowed from Jennifer Widom from Stanford

Movie Example Instance

Movie:

MovielD	Title	Year
1	Star Wars	1977
2	Gone with the Wind	1939
3	The Wizard of Oz	1939
4	Indiana Jones and the Raiders of the Lost Ark	1981

StarsIn:

MovielD	StarID	Role
1	1	Han Solo
4	1	Indiana Jones
2	2	Scarlett O'Hara
3	3	Dorothy Gale

MovieStar:

StarID	Name	Gender
1	Harrison Ford	Male
2	Vivian Leigh	Female
3	Judy Garland	Female

SQL Statements

- **Data Definition Language (DDL)**
 - Statements to define the database schema
- **Data Manipulation Language (DML)**
 - Statements to manipulate the data

CREATE, ALTER and DROP TABLE statements

Basic SELECT Query

Set Operations

Aggregation, GROUP BY and HAVING

Nested Queries

Views

Null Values and Joins

INSERT, DELETE and UPDATE statements

Constraints and Triggers

Data Definition Language (DDL)

- Data Definition Language (DDL) is one of the two main parts to the SQL language.
- DDL statements are used to define the database structure or schema.
 - CREATE - to create objects in the database
 - ALTER - alters the structure of the database
 - DROP - delete objects from the database

Creating Tables in SQL

- **CREATE TABLE** statement creates a new relation, by specifying its **name**, **attributes** and **constraints**.
- The **key**, **entity** and **referential integrity** constraints are specified within the statement after the attributes have been declared.
- The domain constraint is specified for each attribute.
- Data type of an attribute can be specified directly or by declaring a domain (**CREATE DOMAIN**).

Creating Tables in SQL (DDL)

CREATE TABLE <table name>

(<column name> <column type> [<attribute constraint>]
{, <column name> <column type> [<attribute constraint>] }
[<table constraint> {, <table constraint>}])

Interpreting the syntax*

KEYWORD

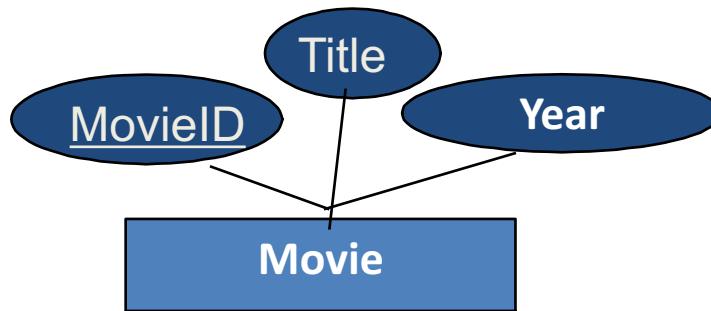
<argument>

[optional]

{multiple}

...|..choice..|...

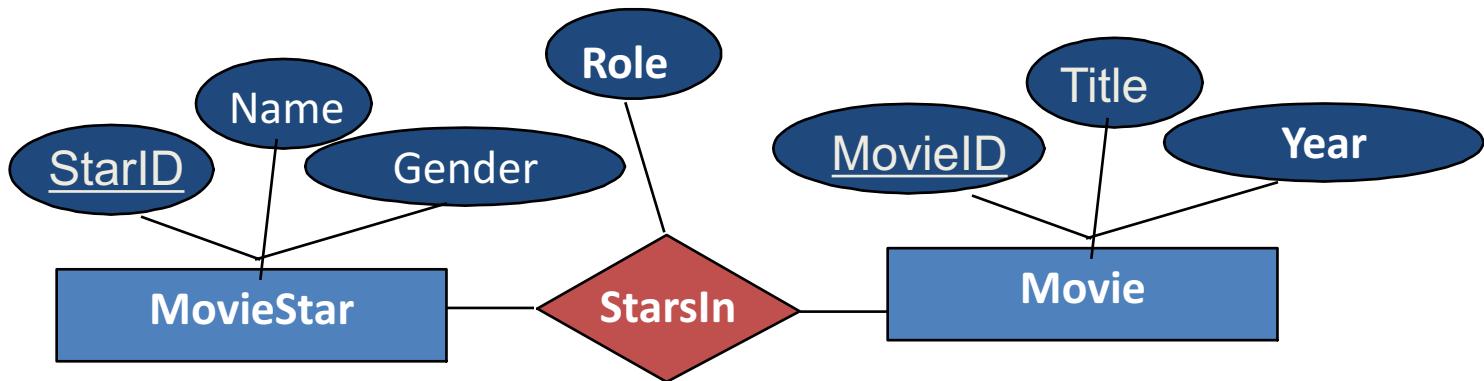
Creating an Entity Table



Movie [MovieID, Title, Year]

```
CREATE TABLE Movie
  (MovieID      INTEGER,
   Title        CHAR(20),
   Year         INTEGER,
   primary key (MovieID))
```

Creating a Relationship Table in SQL



StarsIn[MovieID, StarID, Role]

StarsIn.StarID → **MovieStar**.StarID

StarsIn.MovieID → **Movies**.MovieID

```
CREATE TABLE StarsIn (
    StarID    INTEGER,
    MovieID   INTEGER,
    Role      CHAR(20),
    PRIMARY KEY (StarID, MovieID),
    FOREIGN KEY (StarID) REFERENCES MovieStar,
    FOREIGN KEY (MovieID) REFERENCES Movies)
```

Enforcing Referential Integrity

- MovieID in StarsIn is a foreign key that references Movies
 - StarsIn.MovieID → Movies.MovieID
- What should be done if a *movie tuple* is deleted?
 - Delete all roles that refer to it?
 - Disallow the deletion of the movie?
 - Set MID in WorkOn tuples that refer to it to null?
 - Set MID in WorkOn tuples that refer to it to default value?

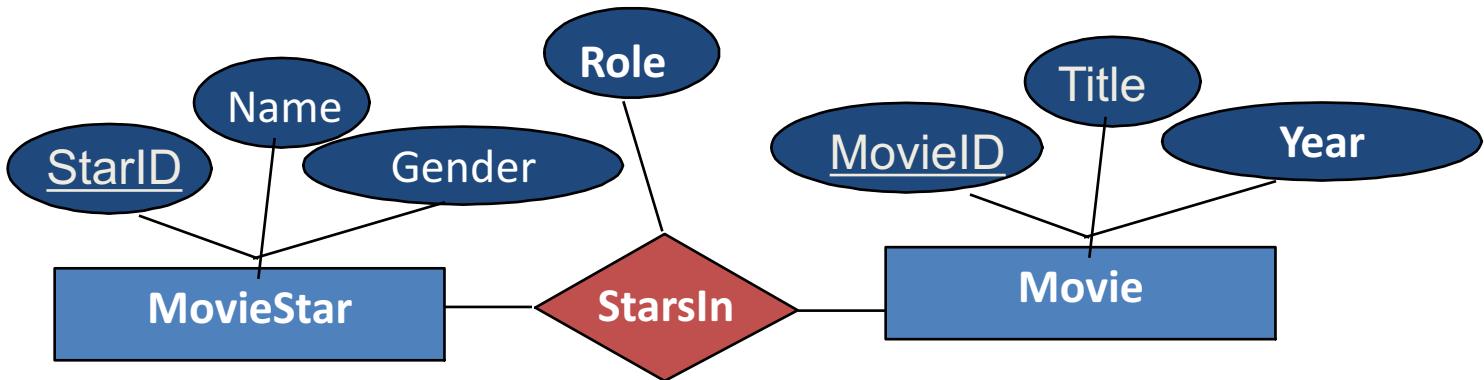
Enforcing Referential Integrity

- A **referential triggered action** clause can be attached to a foreign key constraint, that specifies the action to take if a **referenced tuple** is deleted, or a **referenced primary key** value is modified.
- By default no action is taken and the delete/update is rejected.
- Other actions include the following:

ON DELETE SET NULL | SET DEFAULT | CASCADE

ON UPDATE SET NULL | SET DEFAULT | CASCADE

Creating Tables in SQL (DDL)



```
CREATE TABLE StarsIn (
    StarID    INTEGER,
    MovieID   INTEGER,
    Role      CHAR(20),
    PRIMARY KEY (StarID, MovieID),
    FOREIGN KEY (StarID) REFERENCES MovieStar,
        ON DELETE CASCADE
        ON UPDATE CASCADE
    FOREIGN KEY (MovieID) REFERENCES Movies)
        ON DELETE CASCADE
        ON UPDATE CASCADE
```

Clicker Question

Consider the following table definition.

```
CREATE TABLE BMW ( bid INTEGER, sid INTEGER, ...
    PRIMARY KEY (bid),
    FOREIGN KEY (sid) REFERENCES STUDENTS
        ON DELETE CASCADE);
```

If $\text{bid} = 1000$ and $\text{sid} = 5678$ for a row in Table BMW, choose the best answer

- A. If the row for sid value 5678 in STUDENTS is deleted, then the row with $\text{bid} = 1000$ in BMW is automatically deleted.
- B. If a row with sid value 5678 in BMW is deleted, then the row with $\text{sid}=5678$ in STUDENTS is automatically deleted.
- C. Both of the above.

Clicker Question

Consider the following table definition.

```
CREATE TABLE BMW ( bid INTEGER, sid INTEGER, ...
                    PRIMARY KEY (bid),
                    FOREIGN KEY (sid) REFERENCES STUDENTS
                    ON DELETE CASCADE);
```

If bid = 1000 and sid = 5678 for a row in Table BMW, choose the best answer

- A. If the row for sid value 5678 in STUDENTS is deleted, then the row with bid = 1000 in BMW is automatically deleted. **A is the correct answer**
- B. If a row with sid value 5678 in BMW is deleted, then the row with sid=5678 in STUDENTS is automatically deleted.
- C. Both of the above.

BMW

bid	Sid
1000	5678

Student

sid	name	Address
5678	James	Null

ALTER TABLE

- ALTER TABLE command is used for *schema evolution*, that is the definition of a table created using the CREATE TABLE command, can be changed using the ALTER TABLE command
- Alter table actions include
 - Adding or dropping a column.
 - Changing a column definition.
 - Adding or dropping constraints.

ALTER TABLE Syntax

```
ALTER TABLE <table name>
  ADD <column name> <column type>
    [<attribute constraint>] {, <column name>
      <column type> [<attribute constraint>] }
  | DROP <column name> [CASCADE]
  | ALTER <column name> <column-options>
  | ADD <constraint name> <constraint-options>
  | DROP <constraint name> [CASCADE];
```

DROP TABLE

- **DROP TABLE**
 - Drops all constraints defined on the table including constraints in other tables which reference this table.
 - Deletes all tuples within the table.
 - Removes the table definition from the system catalog.
- **DROP TABLE Syntax**

```
DROP TABLE [IF EXISTS]  
tbl_name [, tbl_name] ...  
[RESTRICT | CASCADE]
```

Basic SELECT Query

Set Operations

Aggregation, GROUP BY and HAVING

Nested Queries

Views

Null Values and Joins

INSERT, DELETE and UPDATE statements

Constraints and Triggers

Data Manipulation Language (DML)

- Data Manipulation Language (DML) is the other main part of the SQL language.
- DML statements are used for managing data within schema objects.
 - SELECT - retrieve data from a database.
 - INSERT - insert data into a table.
 - UPDATE - updates existing data within a table.
 - DELETE – deletes records from a table.

Basic SELECT Query

- In the SELECT statement, users specify what the result of the query should be, and the DBMS decides the operations and order of execution, thus SQL queries are “**Declarative**”.
- The result of a SQL query is a table (relation).
- Note that the SQL SELECT statement has **NO** relationship to the SELECT operation of relational algebra !

Basic SELECT Query

- Selection (WHERE clause)
 - Horizontal scanner to select tuples from given collection of tuples.
- Projection (SELECT clause)
 - Vertically select the attributes of given collection of tuples.
- Join (FROM clause)
 - Combine tuples from different relations for the search purposes.
- Sorting (ORDER clause)
 - Order the resulting tuples according to the given sort key.

SELECT Basic Syntax

```
SELECT <attribute list>
FROM <table list>
[WHERE <condition>] ;
```

- <attribute list> is a list of attribute names whose values are to be retrieved by the query.
- <table list> is a list of relation names required to process the query.
- <condition> is a conditional (Boolean) expression that identifies the tuples to be retrieved by the query.

Projection in SQL

- **Projection** (SELECT clause)
 - Vertically select the attributes of given collection of tuples.

```
SELECT [DISTINCT] (attribute list | * )  
FROM <table list>  
[WHERE <condition>];
```

- **Distinct**: By default, duplicates are not eliminated in SQL relations, which are **bags** or **multisets** and not sets. Use of distinct will eliminate duplicates and enforce set semantics.
- *****: acts as a *wild card*, selecting all of the columns in the table.

Projection Example

- Find the titles of movies.

Movie(MovieID, Title, Year)
StarsIn(MovieID, StarID, Role)
MovieStar(StarID, Name, Gender)

Query

```
SELECT Title  
FROM Movie
```

Results

Title
Star Wars
Gone with the Wind
The Wizard of Oz
Indiana Jones and the Raiders of the Lost Ark

Clicker Question: SQL Projection

- Consider the given table and SQL query.

```
SELECT Score1, Score2  
FROM Scores
```

- Which one of the following tuples is in the result?
 - (1,2)
 - (5,3)
 - (8,6)
 - All are in the answer
 - None are in the answer

Scores			
Team1	Team2	Score1	Score2
Dragons	Tigers	5	3
Carp	Swallows	4	6
Bay Stars	Giants	2	1
Marines	Hawks	5	3
Ham Fighters	Buffaloes	1	6
Lions	Golden Eagles	8	12

Clicker Question: SQL Projection

- Consider the given table and SQL query.

```
SELECT Score1, Score2  
FROM Scores
```

- Which one of the following tuples is in the result?

- A. (1,2)
- B. (5,3) **Correct Answer**
- C. (8,6)
- D. All are in the answer
- E. None are in the answer

Scores		Score1	Score2
Team1	Team2		
Dragons	Tigers	5	3
Carp	Swallows	4	6
Bay Stars	Giants	2	1
Marines	Hawks	5	3
Ham Fighters	Buffaloes	1	6
Lions	Golden Eagles	8	12

clickerprojection.sql

Projection and Duplicates

- Find all the years where a movie was produced.

```
Movie(MovieID, Title, Year)  
StarsIn(MovieID, StarID, Role)  
MovieStar(StarID, Name, Gender)
```

Query

Results

```
SELECT Year  
FROM Movie
```

Year
1977
1939
1939
1981

Query

Results

```
SELECT DISTINCT Year  
FROM Movie
```

Year
1977
1939
1981

Clicker Question on Distinction

- Consider the given table and SQL query.

```
SELECT DISTINCT Team, RunsFor  
FROM Scores
```

Which is true:

- A. 1 appears once
- B. 5 appears twice
- C. 6 appears 4 times
- D. All are true
- E. None are true

Team	Opponent	Runs For	Runs Against
Dragons	Tigers	5	3
Carp	Swallows	4	6
Bay Stars	Giants	2	1
Marines	Hawks	5	3
Ham Fighters	Buffaloes	1	6
Lions	Golden Eagles	8	12
Tigers	Dragons	3	5
Swallows	Carp	6	4
Giants	Bay Stars	1	2
Hawks	Marines	3	5
Buffaloes	Ham Fighters	6	1
Golden Eagles	Lions	12	8

Clicker Question on Distinction

- Consider the given table and SQL query. [clickerdistinction.sql](#)

```
SELECT DISTINCT Team, RunsFor  
FROM Scores
```

Which is true:

- A. 1 appears once
- B. 5 appears twice **Correct**
- C. 6 appears 4 times
- D. All are true
- E. None are true

Team	Opponent	Runs For	Runs Against
Dragons	Tigers	5	3
Carp	Swallows	4	6
Bay Stars	Giants	2	1
Marines	Hawks	5	3
Ham Fighters	Buffaloes	1	6
Lions	Golden Eagles	8	12
Tigers	Dragons	3	5
Swallows	Carp	6	4
Giants	Bay Stars	1	2
Hawks	Marines	3	5
Buffaloes	Ham Fighters	6	1
Golden Eagles	Lions	12	8

Projection and Expressions

- SQL queries can also evaluate expressions and return the value of these expressions together with the projected attributes.
- Expressions use standard arithmetic operators (+, -, *, /) on numeric values or attributes with numeric domains.

Query

Results

```
SELECT Year  
FROM Movie
```

Year
1977
1939
1939
1981

Query

Results

```
SELECT Year+2  
FROM Movie
```

Year
1979
1941
1941
1983

Selection in SQL

- Selection (WHERE clause)
 - Horizontal scanner to select tuples from given collection of tuples.

```
SELECT <attribute list>  
FROM <table list>  
[WHERE join condition and search_condition]
```

<search condition> is a conditional (Boolean) expression that identifies the tuples to be retrieved by the query.

Select Search Example

- Find all of the male stars.

Movie(MovieID, Title, Year)
StarsIn(MovieID, StarID, Role)
MovieStar(StarID, Name, Gender)

Query

Results

```
SELECT *
FROM   MovieStar
WHERE  Gender = "Male"
```

StarID	Name	Gender
1	Harrison Ford	Male

Clicker Question: Selection

- Consider the given table and SQL query.

```
SELECT *
FROM Scores
WHERE RunsFor > 5
```

- Which one of the following tuple is in the result?
 - (Swallows, Carp, 6, 4)
 - (Swallows, Carp, 4)
 - (12)
 - (*)

Team	Opponent	Runs For	Runs Against
Dragons	Tigers	5	3
Carp	Swallows	4	6
Bay Stars	Giants	2	1
Marines	Hawks	5	3
Ham Fighters	Buffaloes	1	6
Lions	Golden Eagles	8	12
Tigers	Dragons	3	5
Swallows	Carp	6	4
Giants	Bay Stars	1	2
Hawks	Marines	3	5
Buffaloes	Ham Fighters	6	1
Golden Eagles	Lions	12	8

Clicker Question: Selection

- Consider the given table and SQL query.

clickerselection.sql

```
SELECT *
FROM Scores
WHERE RunsFor > 5
```

- Which one of the following tuple is in the result?

- A. (Swallows, Carp, 6, 4)
- B. (Swallows, Carp, 4)
- C. (12)
- D. (*)

Team	Opponent	Runs For	Runs Against
Dragons	Tigers	5	3
Carp	Swallows	4	6
Bay Stars	Giants	2	1
Marines	Hawks	5	3
Ham Fighters	Buffaloes	1	6
Lions	Golden Eagles	8	12
Tigers	Dragons	3	5
Swallows	Carp	6	4
Giants	Bay Stars	1	2
Hawks	Marines	3	5
Buffaloes	Ham Fighters	6	1
Golden Eagles	Lions	12	8

Selection Example with Dates

- Find events that have occurred before 1943

events	
name	date
A	1941-05-25
B	1942-11-15
C	1943-12-26
D	1944-10-25

Query

```
SELECT *  
FROM events  
WHERE date < 19430000
```

Results

name	date
A	1941-05-25
B	1942-11-15

Selection & Projection – Together Forever

- What are the names of the female movie stars?

```
SELECT name  
FROM MovieStar  
WHERE Gender = 'female'
```

- What are the titles of movies from prior to 1939?

```
SELECT title  
FROM Movie  
WHERE year < 1939
```

Complex WHERE Conditions

- Find the title of all of the movies that contain “sin”

```
SELECT  *
FROM    movie
Where Title like "%sin%"
```

- LIKE is used for string matching:
 - ‘%’ stands for 0 or more arbitrary characters.
 - ‘_’ stands for any one character.

Complex WHERE Conditions

- Substring Comparisons
 - **LIKE**
 - ... WHERE Address LIKE '%St Lucia%'
 - ... WHERE StrDate LIKE '_ _ / 0 5 / _ _'
 - **IN**
 - ... WHERE LName IN ('Jones', 'Wong', 'Harrison')
 - **IS**
 - ... WHERE DNo IS NULL
- Arithmetic Operators and Functions
 - **+ , - , * , / , date and time functions**, etc.
 - ... WHERE Salary * 2 > 50000
 - ... WHERE Year(Sys_Date - Bdate) > 55
 - **BETWEEN**
 - ... WHERE Salary BETWEEN 10000 AND 30000

Join in SQL

- **Join** (FROM clause)
 - Combine tuples from different relations for the search purposes.

```
SELECT <attribute list>  
FROM <table list of more than one table>  
[WHERE join condition and search_condition]
```

- <join condition> corresponds to a join condition in Relational Algebra.
- Alias for Table names are used to give a table a temporary name to make the query more readable.
 - e.g., FROM StarsIn S

Join in SQL

- Joining R1 and R2 on their shared attribute B:
 - each tuple of R1 is concatenated with every tuple in R2 having the same values on the join attributes.

```
SELECT A, R1.B, C  
FROM R1, R2  
WHERE R1.B = R2.B
```

R ₁	
A	B
1	2
4	5
7	2

R ₂	
B	C
2	3
5	6
2	8

$R_1 \bowtie R_2$



A	B	C
1	2	3
4	5	6
7	2	8
1	2	8
7	2	3

Join Example with Duplication

- Find the ids and names of all movie stars who have been in a movie.

Movie(MovieID, Title, Year)

StarsIn(MovieID, StarID, role)

MovieStar(StarID, Name, Gender)

```
SELECT S.StarID, Name  
FROM StarsIn S, MovieStar MS  
WHERE S.StarID = MS.StarID
```

StarID	Name	Gender
1	Harrison Ford	Male
2	Vivian Leigh	Female
3	Judy Garland	Female

MovieID	StarID	Character
1	1	Han Solo
4	1	Indiana Jones
2	2	Scarlett O'Hara
3	3	Dorothy Gale

The row 1, Harrison Ford will appear twice

```
SELECT DISTINCT S.StarID, Name  
FROM StarsIn S, MovieStar MS  
WHERE S.StarID = MS.StarID
```

Join Example

- Find the ids, names and characters of all movie stars who have been in the movie with MovieID 1

```
Movie(MovieID, Title, Year)  
StarsIn(MovieID, StarID, Role)  
MovieStar(StarID, Name, Gender)
```

Query

```
SELECT S.StarID, Name, Role  
FROM StarsIn S, MovieStar MS  
WHERE S.StarID = MS.StarID and  
S.MovieID = 1
```

Results

StarID	Name	Role
1	Harrison Ford	Han Solo

Join Example - Complex Conditions

- Find the ids, names and characters of all movie stars who have been in the movie titled ‘Gone with the Wind’.

Movie(MovieID, Title, Year)
StarsIn(MovieID, StarID, Role)
MovieStar(StarID, Name, Gender)

Query

```
SELECT S.StarID, Name, Role, M.title  
FROM StarsIn S, MovieStar MS, Movie M  
WHERE S.StarID = MS.StarID and  
S.MovieID = M.MovieID and  
M.title like "Gone with the Wind"
```

Results

StarID	Name	Role
2	Scarlett O'Hara	Scarlett O'Hara

Clicker Question: Joins

- Consider R :

a	b
0	0
0	1
1	0
1	1

S:

a	b
0	0
0	1
1	0
1	1

T:

a	b
0	0
0	1
1	0
1	1

```
SELECT R.a, R.b, S.b, T.b  
FROM R, S, T  
WHERE R.b = S.a AND S.b <> T.b  (note: <> == 'not equals')
```

Compute the results. Which of the following are true:

- (0,1,1,0) appears twice.
- (1,1,0,1) does not appear.
- (1,1,1,0) appears once.
- All are true
- None are true

Clicker Question: Joins

- Consider R :

a	b
0	0
0	1
1	0
1	1

S:

a	b
0	0
0	1
1	0
1	1

T:

a	b
0	0
0	1
1	0
1	1

```
SELECT R.a, R.b, S.b, T.b
FROM R, S, T
WHERE R.b = S.a AND S.b <> T.b  (note: <> == 'not equals')
```

Compute the results. Which of the following are true:

A. (0,1,1,0) appears twice.

True R(0,1) S(1,1), T(0,0)&R(0,1), S(1,1), T(1,0)

B. (1,1,0,1) does not appear.

False: R(1,1), S(1,0), T(0,1)

C. (1,1,1,0) appears once.

False: like A but use R(1, 1)

D. All are true

E. None are true

Renaming Attributes

- SQL allows renaming relations and attributes using the **as** clause:
old-name as new-name
- Example: Find the title of movies and all the characters in them, and rename “Role” to “Role1”.

```
SELECT Title, Role AS Role1  
FROM StarsIn S, Movie M  
WHERE M.MovieID = S.MovieID
```

Try select *; does not remove duplicate columns

Sorting in SQL

- **Sorting** (ORDER clause)
 - Order the resulting tuples according to the given sort key.

```
SELECT [DISTINCT] (attribute/expression list | * )
FROM <table list>
[WHERE [join condition and] search_condition]
[ORDER BY column_name [ASC|DESC] {, column-name [ASC|DESC]}];
```

Order is specified by:

- **asc** for ascending order (default)
- **desc** for descending order
- E.g. **order by Name desc**

Ordering of Tuples

- List in alphabetic order the names of actors who were in a movie in 1939.

Movie(MovieID, Title, Year)
StarsIn(MovieID, StarID, Role)
MovieStar(StarID, Name, Gender)

Query

```
SELECT distinct Name  
FROM Movie M, StarsIn S, MovieStar MS  
WHERE M.MovieID = S.MovieID and  
S.StarID = MS.StarID and year = 1939  
ORDER BY Name
```

Results

Name
Vivian Leigh
Judy Garland

Clicker question: sorting

- Consider the following query:

```
SELECT a, b, c  
FROM R  
ORDER BY c DESC, b ASC;
```
- What condition must a tuple t satisfy so that t **necessarily precedes** the tuple $(5,5,5)$? Identify one such tuple from the list below.
 - A. $(3,6,3)$
 - B. $(1,5,5)$
 - C. $(5,5,6)$
 - D. All of the above
 - E. None of the above

Clicker question: sorting

- Consider the following query:

```
SELECT a, b, c  
FROM R  
ORDER BY c DESC, b ASC;
```

- What condition must a tuple t satisfy so that t **necessarily precedes** the tuple $(5,5,5)$? Identify one such tuple from the list below.

A. $(3,6,3)$

3 < 5

B. $(1,5,5)$

Not specified

C. $(5,5,6)$

Correct

D. All of the above

E. None of the above

clickerorder.sql and
clickerorder2.sql produce different
ordering for 7,5,5 vs. 1,5,5

Conceptual Procedural Evaluation Strategy

1. Compute the cross-product of *relation-list*.
2. Discard resulting tuples if they fail *qualifications*.
3. Delete attributes that are not in *target-list*.
4. If DISTINCT is specified, eliminate duplicate rows.
5. If ORDER BY is specified, sort the results.

CREATE, ALTER and DROP TABLE statements

Basic SELECT Query

Set Operations

Aggregation, GROUP BY and HAVING

Nested Queries

Views

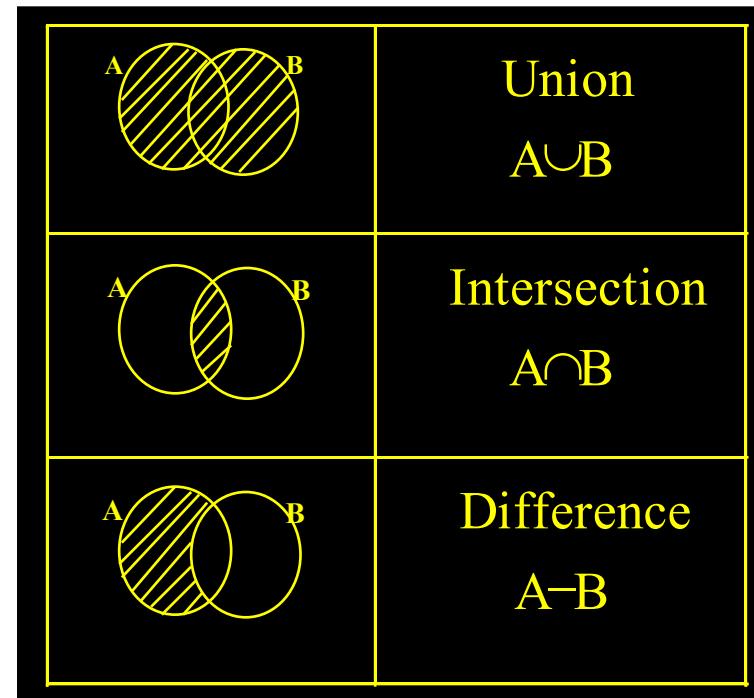
Null Values and Joins

INSERT, DELETE and UPDATE statements

Constraints and Triggers

Basic Set Operators

- Relation is a *set* of tuples (no duplicates).
- Set theory, and hence elementary set operators also apply to relations
 - UNION
 - INTERSECTION
 - DIFFERENCE



Set Operations

- **Union, intersect, and except** correspond to the relational algebra operations \cup , \cap , $-$.
- Each automatically eliminates duplicates;
 - To retain all duplicates use the corresponding multiset versions:
union all, intersect all and **except all**.

- Suppose a tuple occurs m times in r and n times in s , then, it occurs:
 - $m + n$ times in $r \text{ union all } s$
 - $\min(m, n)$ times in $r \text{ intersect all } s$
 - $\max(0, m - n)$ times in $r \text{ except all } s$

Union Compatibility

Two relations $R1(A1, A2, \dots, An)$ and $R2(B1, B2, \dots, Bn)$ are *union compatible* iff:

- They have the same degree n, (number of columns).
- Their columns have corresponding domains, i.e $\text{dom}(Ai) = \text{dom}(Bi)$ for $1 \leq i \leq n$
- Note that although domains need to correspond they do not have to have the same name.

Set Operations: Union

- **UNION:** Produces a relation that includes all tuples that appear only in R1, or only in R2, or in both R1 and R2.
 - Duplicate Tuples are eliminated if UNION ALL is not used.
 - R1 and R2 must be union compatible.

```
SELECT ...  
UNION [ALL] SELECT ...  
[UNION [ALL] SELECT ...]
```

Union Example

- Find IDs of MovieStars who've been in a movie in 1944
or 1974.

```
Movie(MovieID, Title, Year)  
StarsIn(MovieID, StarID, Role)  
MovieStar(StarID, Name, Gender)
```

```
SELECT StarID  
FROM Movie M, StarsIn S  
WHERE M.MovieID=S.MovieID AND  
( year = 1944 OR year = 1974)
```

```
SELECT StarID  
FROM Movie M, StarsIn S  
WHERE M.MovieID = S.MovieID  
AND year = 1944  
UNION  
SELECT StarID  
FROM Movie M, StarsIn S  
WHERE M.MovieID = S.MovieID  
AND year = 1974
```

- Are the queries the same?

Intersection in SQL

- **Intersection:** Produces a relation that includes the tuples that appear in both R1 and R2.
 - **Duplicate** Tuples are eliminated if INTERSECT ALL is not used.
 - R1 and R2 must be union compatible.

```
SELECT ...
INTERSECT [ALL] SELECT ...
[INTERSECT [ALL] SELECT ...]
```

Intersect Example

- Find IDs of stars who have been in a movie in 1944 and 1974.

```
Movie(MovieID, Title, Year)  
StarsIn(MovieID, StarID, Role)  
MovieStar(StarID, Name, Gender)
```

```
SELECT StarID  
FROM Movie M, StarsIn S  
WHERE M.MovieID=S.MovieID AND  
( year = 1944 AND year = 1974)
```

```
SELECT StarID  
FROM Movie M, StarsIn S  
WHERE M.MovieID = S.MovieID  
AND year = 1944  
INTERSECT  
SELECT StarID  
FROM Movie M, StarsIn S  
WHERE M.MovieID = S.MovieID  
AND year = 1974
```

INTERSECT is part of the SQL standard, but is not implemented in MySQL.

Rewriting INTERSECT with Joins

- Example: Find IDs of stars who have been in a movie in 1944 *and* 1974 without using **INTERSECT**.

Movie(MovieID, Title, Year)
StarsIn(MovieID, StarID, Role)
MovieStar(StarID, Name, Gender)

```
SELECT distinct S1.StarID
FROM   Movie M1, StarsIn S1,
       Movie M2, StarsIn S2
WHERE
      M1.MovieID = S1.MovieID AND M1.year = 1944 AND
      M2.MovieID = S2.MovieID AND M2.year = 1974 AND
      S2.StarID = S1.StarID
```

Rewriting INTERSECT with Joins

Movie(MovieID, Title, Year)
StarsIn(MovieID, StarID, Role)
MovieStar(StarID, Name, Gender)

```
SELECT distinct S1.StarID  
FROM Movie M1, StarsIn S1  
WHERE M1.MovieID = S1.MovieID  
AND M1.year = 1944
```

StarID

1043
1044
1045
1046

```
SELECT distinct S2.StarID  
FROM Movie M2, StarsIn S2  
WHERE M2.MovieID = S2.MovieID  
AND M2.year = 1974
```

StarID

1149
1150
1151
1045

```
SELECT distinct S1.StarID  
FROM Movie M1, StarsIn S1,  
Movie M2, StarsIn S2  
WHERE  
M1.MovieID = S1.MovieID AND M1.year = 1944 AND  
M2.MovieID = S2.MovieID AND M2.year = 1974 AND  
S2.StarID = S1.StarID
```

Difference in SQL

- EXCEPT(also referred to as MINUS) Produces a relation that includes all the tuples that appear in R1, but do not appear in R2.
 - R1 and R2 must be union compatible.

```
SELECT ...  
EXCEPT [ALL] SELECT ...  
[EXCEPT [ALL] SELECT ...]
```

EXCEPT Example

- Find IDs of stars who have been in a movie in 1944 but not in 1974.

```
Movie(MovieID, Title, Year)
StarsIn(MovieID, StarID, Role)
MovieStar(StarID, Name, Gender)
```

```
SELECT StarID
FROM Movie M, StarsIn S
WHERE M.MovieID = S.MovieID AND
year = 1944
Except
SELECT StarID
FROM Movie M, StarsIn S
WHERE M.MovieID = S.MovieID AND
year = 1974
```

EXCEPT is part of the SQL standard, but is not implemented in MySQL.

EXCEPT queries can be implemented using nested queries

Properties of Set Operators

$A \cup B$	commutative	$A \cup B = B \cup A$
	associative	$(A \cup B) \cup C = A \cup (B \cup C)$
$A \cap B$	commutative	$A \cap B = B \cap A$
	associative	$(A \cap B) \cap C = A \cap (B \cap C)$
$A - B$	not commutative	$A - B \neq B - A$
	not associative	$(A - B) - C \neq A - (B - C)$

CREATE, ALTER and DROP TABLE statements

Basic SELECT Query

Set Operations

Aggregation, GROUP BY and HAVING

Nested Queries

Views

Null Values and Joins

INSERT, DELETE and UPDATE statements

Constraints and Triggers

Aggregation in SQL

- Aggregates are functions that produce summary values.

```
SELECT [DISTINCT] (attribute / exprsn / aggregation-function list | * )
```

```
FROM <table list>
```

```
[WHERE [join condition and] search_condition]
```

```
[ORDER BY column_name [ASC|DESC] {, column-name [ASC|DESC]}];
```

- The **aggregation-function** list may include:

- **SUM/ AVG ([DISTINCT] expression)**: Calculates the sum/ average of a set of *numeric* values
- **COUNT ([DISTINCT] expression)**: Counts the number of tuples that the query returns
- **COUNT(*)**
- **MAX/MIN(expression)**: Returns the maximum (minimum) value from a set of values which have a *total ordering*. Note that the domain of values can be non-numeric.

Aggregate Operators Examples

College(cName, state, enrollment)

Student(sID, sName, GPA, sizeHS)

Apply(sID, cName, major, decision)

students

```
SELECT COUNT(*)  
FROM Student
```

Finding average GPA of students
from high schools with less than 500
students

```
SELECT AVG (GPA)  
FROM Student  
WHERE sizeHS<500
```

Aggregation Examples

College(cName, state, enrollment)
Student(sID, sName, GPA, sizeHS)
Apply(sID, cName, major, decision)

- Find the minimum GPA.

```
SELECT min(GPA)
FROM Student
```

- Find how many students have applied to ‘Stanford’.

```
SELECT count(distinct sID)
FROM Apply
where cname like 'Stanford'
```

Note: want distinct for when Students apply to more than one major at Stanford

GROUP BY and HAVING

- Divide tuples into groups and apply aggregate operations to each group.
- Example: *Find the age of the youngest student for each major.*

For $i = \text{'Computer Science'}$,
 $\text{'Civil Engineering'}$...

```
SELECT MIN (age)
FROM Student
WHERE major = i
```

■ Problem:

We don't know how many majors exist, not to mention this is not good practice

GROUP BY Syntax

- Aggregation functions can also be applied to groups of rows within a table. The GROUP BY clauses provides this functionality.

```
SELECT [DISTINCT] (attribute / expression / aggregation-function list | * )  
FROM <table list>  
[WHERE [join condition and]      search_condition]  
[GROUP BY grouping attributes]  
[ORDER BY column_name [ASC|DESC] {, column-name [ASC|DESC]}];
```

- When GROUP BY is used in an SQL statement, any attribute appeared in SELECT Clause must also appeared in an aggregation function or in GROUP BY clause.

Grouping Examples

- Find the age of the youngest student who is at least 19, for each major.

```
SELECT      major, MIN(age)
FROM        Student
WHERE       age >= 19
GROUP BY    major
```

Snum	Major	Age
115987938	Computer Science	20
112348546	Computer Science	19
280158572	Animal Science	18
351565322	Accounting	19
556784565	Civil Engineering	21
...

Major	Age
Computer Science	19
Accounting	19
Civil Engineering	21
...	...

No Animal Science

Conditions on Groups

- Conditions can be imposed on the selection of groups to be included in the query result.
- The HAVING clause (following the GROUP BY clause) is used to specify these conditions, similar to the WHERE clause.

```
SELECT [DISTINCT] (attribute / expression / aggregation-function list | * )  
FROM <table list>  
[WHERE [join condition and] search_condition]  
[GROUP BY grouping attributes]  
[HAVING <group condition>]  
[ORDER BY column_name [ASC|DESC] {, column-name [ASC|DESC]}];
```

- Unlike the WHERE clause, the HAVING clause can also include aggregates.

Grouping Examples with Having

- Find the age of the youngest student who is at least 19, for each major with at least 2 such students.

```
SELECT      major, MIN(age)
FROM        Student
WHERE       age >= 19
GROUP BY    major
HAVING     COUNT(*) > 1
```

Major	Age
Computer Science	19
Accounting	19
Civil Engineering	21
...	...



Major	Age
Computer Science	19

Snum	Major	Age
115987938	Computer Science	20
112348546	Computer Science	19
280158572	Animal Science	18
351565322	Accounting	19
556784565	Civil Engineering	21
...

- Can we select age instead of MIN(age)?
- Can we select snum to be returned?
- Can we select major to be returned?

GROUP BY and HAVING (cont)

```
SELECT      [DISTINCT] target-list
FROM        relation-list
WHERE       qualification
GROUP BY    grouping-list
HAVING      group-qualification
ORDER BY    target-list
```

- The *target-list* contains
 - (i) attribute names
 - (ii) terms with aggregate operations (e.g., MIN (*S.age*)).
- Attributes in (i) must also be in *grouping-list*.
 - each answer tuple corresponds to a *group*,
 - *group* = a set of tuples with same value for all attributes in *grouping-list*
 - selected attributes must have a single value per group.
- Attributes in *group-qualification* are either in *grouping-list* or are arguments to an aggregate operator.

Conceptual Evaluation of a Query

1. compute the cross-product of *relation-list*.
2. keep only tuples that satisfy *qualification*.
3. partition the remaining tuples into groups by where attributes in *grouping-list*.
4. keep only the groups that satisfy *group-qualification* (expressions in *group-qualification* must have a single value per group!).
5. delete fields that are not in *target-list*.
6. generate one answer tuple per qualifying group.

Clicker Question on Grouping

- Compute the result of the query:

```
SELECT a1.x, a2.y, COUNT(*)  
FROM Arc a1, Arc a2  
WHERE a1.y = a2.x  
GROUP BY a1.x, a2.y
```

x	y
1	2
1	2
2	3
3	4
3	4
4	1
4	1
4	1
4	2

Which of the following is in the result?

- (1,3,2)
- A. (4,2,6)
- B. (4,3,1)
- C. All of the above
- D. None of the above

Tip: You can think of Arc as being a flight, and the query as asking for how many ways you can take each 2 hop plane trip

Clicker Question on Grouping

- Compute the result of the query:

```
SELECT a1.x, a2.y, COUNT(*)
FROM Arc a1, Arc a2
WHERE a1.y = a2.x
GROUP BY a1.x, a2.y
```

x	y	COUNT(*)
1	3	2
2	4	2
3	1	6
3	2	2
4	2	6
4	3	1

Which is in the result?

- (1,3,2) (1,2)(2,3), (1,2)(2,3)
- A. (4,2,6) 3 ways to do (4,1) and two ways to do (1,2)
- B. (4,3,1) (4,2)(2,3)
- C. All of the above Correct
- D. None of the above

Tip: You can think of Arc as being a flight, and the query as asking for how many ways you can take each 2 hop plane trip

Clicker Question: Having

Suppose we have a relation with schema R(A, B, C, D, E). If we issue a query of the form:

```
SELECT ...
FROM R
WHERE ...
GROUP BY B, E
HAVING ???
```

What terms can appear in the HAVING condition (represented by **???** in the above query)? Identify, in the list below, the term that **CANNOT** appear.

- A. A
- B. B
- C. Count (B)
- D. All can appear
- E. None can appear

Clicker Question: Having

Suppose we have a relation with schema R(A, B, C, D, E). If we issue a query of the form:

```
SELECT ...
FROM R
WHERE ...
GROUP BY B, E
HAVING ???
```

Any aggregated term can appear in HAVING clause. An attribute not in the GROUP-BY list cannot be unaggregated in the HAVING clause. Thus, B or E may appear unaggregated, and all five attributes can appear in an aggregation. However, A, C, or D cannot appear alone.

What terms can appear in the HAVING condition (represented by ??? in the above query)? Identify, in the list below, the term that **CANNOT** appear.

- A. A **Correct. A cannot appear unaggregated**
- B. B
- C. Count (B)
- D. All can appear
- E. None can appear

Grouping Examples

- Find the age of the youngest student with age > 18, for each major with at least 2 students (of age > 18).

Student (Snum, Major, Age)

```
SELECT S.major, MIN(S.age)
FROM Student S
WHERE S.age > 18
GROUP BY S.major
HAVING count(*) >1
```

CREATE, ALTER and DROP TABLE statements

Basic SELECT Query

Set Operations

Aggregation, GROUP BY and HAVING

Nested Queries

Views

Null Values and Joins

INSERT, DELETE and UPDATE statements

Constraints and Triggers

Motivating Example for Nested Queries

- Find ids and names of stars who have been in movie with ID 28:

```
SELECT Distinct M.StarID, name  
FROM MovieStar M, StarsIn S  
WHERE M.StarID = S.starID AND S.MovieID = 28;
```

- Find ids and names of stars who have not been in movie with ID 28:

```
SELECT Distinct M.StarID, name  
FROM MovieStar M, StarsIn S  
WHERE M.StarID = S.starID AND S.MovieID <> 28;
```

Nested Queries in SQL

- Concept of nested sub-queries
- Correlated and non-correlated variants
- Sub-query operators
- Sub-queries vs. set operations and joins
- Division in SQL

Nested SQL Queries

- A nested query (often termed sub-query) is a query that appears within another query.
 - Inside the WHERE clause of another SELECT statement.
 - Inside an INSERT, UPDATE or DELETE statement.
 - Nesting can occur at multiple levels.
- Nested queries are useful for expressing queries where data must be fetched and used in a comparison condition.

Syntax of Nested SQL Queries

```
SELECT ... FROM ... WHERE  
    {expression {[NOT] IN |  
    comparison-operator [ANY|ALL]}  
    | [NOT] EXISTS}  
(SELECT ... FROM ... WHERE ...);
```

Outer Query

Nested/Sub-Query

- Other query search conditions (including joins) can also appear in the outer query WHERE clause, either before or after the inner query

Nested Queries in SQL

- Concept of nested sub-queries
- Correlated and non-correlated variants
- Sub-query operators
- Sub-queries vs. set operations and joins
- Division in SQL

Non-correlated Nested Queries

- Non-correlated Nested Queries
 - Results are returned from an inner query to an outer clause, that is sub-queries are evaluated from the “inside out”.
 - The outer query takes an action based on the results of the inner query.

Non-correlated Nested Queries Example

- Find ids and names of stars who have been in movie with ID 28:

```
SELECT M.StarID, M.Name  
FROM MovieStar M  
WHERE M.StarID IN (SELECT S.StarID  
                    FROM StarsIn S  
                    WHERE MovieID=28)
```

NOT IN

- In this example, the inner query does not depend on the outer query so it could be computed just once.
- Think of this as a function that has no parameters.

```
SELECT S.StarID  
FROM StarsIn S  
WHERE MovieID=28
```

StarID
1026
1027

```
SELECT M.StarID, M.Name  
FROM MovieStar M  
WHERE M.StarID IN  
(1026,1027)
```

Correlated Nested Queries

- Correlated Nested Queries
 - Correlated subqueries have **conditions in their WHERE clause** that references some attribute of a relation declared in the outer query.
 - The outer SQL statement provides the values for the inner subquery to use in its evaluation.
 - The subquery is evaluated once for each (combination of) tuple in the outer query.

Correlated Nested Queries Example

- For each college, check if there is another college in the same state

```
SELECT cName, state  
FROM College C1  
WHERE exists (SELECT *  
              FROM College C2  
              WHERE C2.state = C1.state AND  
                    C2.cName <> C1.cName);
```

Think of this as passing parameters

Outer Query

cName	state
Stanford	CA
Berkeley	CA
MIT	MA
Cornell	NY

Results

cName	state
Stanford	CA

Correlated Nested Queries Example

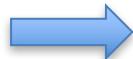
- For each college, check if there is another college in the same state

```
SELECT cName, state  
FROM College C1  
WHERE exists (SELECT *  
              FROM College C2  
              WHERE C2.state = C1.state AND  
                    C2.cName <> C1.cName);
```

Think of this as passing parameters

Outer Query

cName	state
Stanford	CA
Berkeley	CA
MIT	MA
Cornell	NY



Results

cName	state
Stanford	CA
Berkeley	CA

Correlated Nested Queries Example

- For each college, check if there is another college in the same state

```
SELECT cName, state  
FROM College C1  
WHERE exists (SELECT *  
              FROM College C2  
              WHERE C2.state = C1.state AND  
                    C2.cName <> C1.cName);
```

Think of this as passing parameters

Outer Query

cName	state
Stanford	CA
Berkeley	CA
MIT	MA
Cornell	NY

Results

cName	state
Stanford	CA
Berkeley	CA



Correlated Nested Queries Example

- For each college, check if there is another college in the same state

```
SELECT cName, state  
FROM College C1  
WHERE exists (SELECT *  
              FROM College C2  
              WHERE C2.state = C1.state AND  
                    C2.cName <> C1.cName);
```

Think of this as passing parameters

Outer Query

cName	state
Stanford	CA
Berkeley	CA
MIT	MA
Cornell	NY

Results

cName	state
Stanford	CA
Berkeley	CA



Nested Queries in SQL

- Concept of nested sub-queries
- Correlated and non-correlated variants
- **Sub-query operators**
- Sub-queries vs. set operations and joins
- Division in SQL

Sub-query Operators

- Sub-queries that return a set
 - expression {[NOT] IN (*sub-query*)}
 - expression comp-op [ANY|ALL] (*sub-query*)
- Subqueries that return a single value
 - expression comp-op (*sub-query*)

Sub-query returning a Set

Expression and attribute list in sub-query SELECT clause must have same domain

- expression {[NOT] IN (*sub-query*)}
 - expression is checked for **membership** in the set (of tuples) returned by sub-query.
- expression comp-op [ANY|ALL] (*sub-query*)
 - expression is **compared** with the set (of tuples) returned by the sub-query
 - ANY: Evaluates to true if one comparison is true
 - ALL: Evaluates to true if all comparisons are true

IN/NOT IN Operator Example

- Find ids and names of stars who have been in movie with ID 28:

```
SELECT M.StarID, M.Name  
FROM MovieStar M  
WHERE M.StarID IN (SELECT S.StarID  
                    FROM StarsIn S  
                    WHERE MovieID=28)
```

NOT IN

- In this example, the inner query does not depend on the outer query so it could be computed just once.
- Think of this as a function that has no parameters.

```
SELECT S.StarID  
FROM StarsIn S  
WHERE MovieID=28
```

StarID
1026
1027

```
SELECT M.StarID, M.Name  
FROM MovieStar M  
WHERE M.StarID IN  
(1026,1027)
```

ANY/ ALL Operator Example

- Also available: **op ANY, op ALL**, where **op** is one of: $>$, $<$, $=$, \leq , \geq , \neq
- Find movies made after “Fargo”

outer query

```
SELECT *
FROM Movie
WHERE year > ANY (SELECT year
                     FROM Movie
                     WHERE Title = 'Fargo')
```

Just returning
one column

inner query

if using **ALL**, it returns the tuples that year is greater than
all the result of inner query

- Assuming we have multiple movies named Fargo, how would the use of ALL vs. ANY affect the result?

Equivalence of IN and = ANY

- Find ids and names of stars who have been in movie with ID 28:

```
SELECT M.StarID, M.Name  
FROM MovieStar M  
WHERE M.StarID IN (SELECT S.StarID  
                      FROM StarsIn S  
                      WHERE MovieID=28)
```

```
SELECT M.StarID, M.Name  
FROM MovieStar M  
WHERE M.StarID = ANY (SELECT S.StarID  
                      FROM StarsIn S  
                      WHERE MovieID=28)
```

Equivalence of Not IN and \neq ALL

- Find ids and names of stars who have NOT been in movie with ID 28:

```
SELECT M.StarID, M.Name  
FROM MovieStar M  
WHERE M.StarID NOT IN (SELECT S.StarID  
                          FROM StarsIn S  
                          WHERE MovieID=28)
```

```
SELECT M.StarID, M.Name  
FROM MovieStar M  
WHERE M.StarID <> ALL (SELECT S.StarID  
                          FROM StarsIn S  
                          WHERE MovieID=28)
```

Non-Equivalence of Not IN and \neq ANY

- If a sub-query returns: {\$30K, \$32K, \$37K}
- NOT IN means
 - NOT=\$30K AND NOT=\$32K AND NOT=\$37K
- \neq ANY means
 - NOT=\$30K OR NOT=\$32K OR NOT=\$37K

\neq ANY will be true for any value in this example.

Clicker Nested Question

Consider the following table and SQL query:

```
SELECT Team, Day  
FROM Scores S1  
WHERE Runs <= ALL  
(SELECT Runs  
FROM Scores S2  
WHERE S1.Day = S2.Day )
```

Which of the following is in the result:

- A. (Carp, Sun)
- B. (Bay Stars, Sun)
- C. (Swallows, Mon)
- D. All of the above
- E. None of the above

Team	Day	Opponent	Runs
Dragons	Sun	Swallows	4
Tigers	Sun	Bay Stars	9
Carp	Sun	Giants	2
Swallows	Sun	Dragons	7
Bay Stars	Sun	Tigers	2
Giants	Sun	Carp	4
Dragons	Mon	Carp	6
Tigers	Mon	Bay Stars	5
Carp	Mon	Dragons	3
Swallows	Mon	Giants	0
Bay Stars	Mon	Tigers	7
Giants	Mon	Swallows	5

Nested Grouping Examples

student (Snum, Major, Age)

- Find the age of the youngest student with age > 18, for each major with at least 2 students(of age > 18).

```
SELECT S.major, MIN(S.age)
FROM student S
WHERE S.age > 18
GROUP BY S.major
HAVING count(*) >1
```

- Find the age of the youngest student with age > 18, for each major for which their average age is higher than the average age of all students across all majors.

```
SELECT S.major, MIN(S.age), AVG(age)
FROM student S
WHERE S.age>18
GROUP BY S.major
HAVING AVG(age) > (SELECT AVG(age) FROM S)
```

Nested Grouping Examples

student (Snum, Major, Age)

- Find the age of the youngest student with age > 18, for each major with at least 2 students(of any age).

```
SELECT S.major, MIN(S.age)
FROM student S
WHERE S.age>18
GROUP BY S.major
HAVING 1<(SELECT COUNT(*) FROM S2 WHERE S.major = S2.major)
```

Sub-query Returning a Set

- The select list of an inner sub-query introduced with a comparison operator (and ANY/ALL) or IN can include only **one expression or column name**.
- The expression you name in the WHERE clause of the outer statement must be **join compatible** with the column you name in the sub-query select list.

```
SELECT M.StarID, M.Name  
FROM MovieStar M  
WHERE M.StarID IN (SELECT S.StarID  
                    FROM StarsIn S  
                    WHERE MovieID=28)
```

Sub-query Returning a Value

Expression and attribute list in sub-query SELECT clause must have same domain.

- **expression comp-op** (sub-query)
 - expression is **compared** with the *value* returned by the sub-query.
 - The sub-query must evaluate to a single value otherwise an error will occur.

```
SELECT Title  
FROM Movie  
Where year = (SELECT max(year)  
                FROM Movie)
```

Using the Exists Function

- The EXISTS function tests for the existence or nonexistence of data that meet the criteria of the sub-query

SELECT ... FROM ...

WHERE [NOT] EXISTS (*sub-query*)

- Sub-queries are used with EXISTS and NOT EXISTS are always correlated
 - WHERE EXISTS (*sub-query*) evaluates to true if the result of the correlated sub-query is a non-empty set, i.e. **contains 1 or more tuples.**
 - WHERE NOT EXISTS (*sub-query*) evaluates to true if the result of the correlated sub-query returns **an empty set**, i.e. zero tuples.

Exists NOT EXISTS Example

- Find movies that were the only movie of the year.

```
SELECT *
From Movie M1
where NOT EXISTS
(Select *
 FROM Movie M2
 Where M1.movieid <> M2.movieID and M1.year = M2.year)
```

- Find movies that were not the only movie of the year.

```
SELECT *
From Movie M1
where EXISTS
(Select *
 FROM Movie M2
 Where M1.movieid <> M2.movieID and M1.year = M2.year)
```

Sub-query using ORDER BY

- Sub-queries cannot include the **ORDER BY clause**.
The optional DISTINCT keyword may effectively order the results of a sub-query, since most systems eliminate duplicates by first ordering the results

Nested Queries in SQL

- Concept of nested sub-queries
- Correlated and non-correlated variants
- Sub-query operators
- Sub-queries vs. set operations and joins
- Division in SQL

Sub-queries vs. Set Operations and Joins

- Find IDs of stars who have been in movies in 1944 and 1974.

```
SELECT StarID  
FROM Movie M, StarsIn S  
WHERE M.MovieID = S.MovieID AND  
year = 1944
```

INTERSECT

```
SELECT StarID  
FROM Movie M, StarsIn S  
WHERE M.MovieID = S.MovieID AND  
year = 1974
```

```
SELECT distinct S1.StarID  
FROM Movie M1, StarsIn S1,  
Movie M2, StarsIn S2  
WHERE  
M1.MovieID = S1.MovieID AND M1.year = 1944 AND  
M2.MovieID = S2.MovieID AND M2.year = 1974 AND  
S2.StarID = S1.StarID
```

```
SELECT S.StarID  
FROM Movie M, StarsIn S  
WHERE M.MovieID = S.MovieID AND M.year = 1944 AND  
S.StarID IN (SELECT S2.StarID  
FROM Movie M2, StarsIn S2  
WHERE M2.MovieID = S2.MovieID AND M2.year = 1974)
```

Sub-queries vs. Set Operations and Joins

- Find IDs of stars who have been in movies in 1944 and 1974.

```
SELECT StarID  
FROM Movie M, StarsIn S  
WHERE M.MovieID = S.MovieID  
AND year = 1944  
UNION  
SELECT StarID  
FROM Movie M, StarsIn S  
WHERE M.MovieID = S.MovieID  
AND year = 1974
```

```
SELECT StarID  
FROM Movie M, StarsIn S  
WHERE M.MovieID=S.MovieID AND ( year = 1944 OR year = 1974)
```

```
SELECT S.StarID  
FROM Movie M, StarsIn S  
WHERE M.MovieID = S.MovieID AND M.year = 1944 OR  
S.StarID IN (SELECT S2.StarID  
              FROM Movie M2, StarsIn S2  
              WHERE M2.MovieID = S2.MovieID AND M2.year = 1974)
```

Sub-queries vs. Set Operations in Difference

- Find IDs of stars who have been in movies in 1944 and not in 1974.

```
SELECT StarID  
FROM Movie M, StarsIn S  
WHERE M.MovieID = S.MovieID  
AND year = 1944
```

Except

```
SELECT StarID  
FROM Movie M, StarsIn S  
WHERE M.MovieID = S.MovieID  
AND year = 1974
```

```
SELECT S.StarID  
FROM Movie M, StarsIn S  
WHERE M.MovieID = S.MovieID AND M.year = 1944 AND  
S.StarID NOT IN (SELECT S2.StarID  
FROM Movie M2, StarsIn S2  
WHERE M2.MovieID = S2.MovieID AND M2.year = 1974)
```

Nested Queries Example

- Find IDs and names of students applying to CS (using both join and nested queries).

College(cName, state, enrollment)
Student(sID, sName, GPA, sizeHS)
Apply(sID, cName, major, decision)

Join

```
SELECT DISTINCT sName, S.sID  
FROM Apply A, Student S  
WHERE A.major = 'CS' and A.sID = S.sID
```

Nested queries

```
SELECT sName, sID  
FROM Student  
WHERE sID in (SELECT sID  
              FROM Apply  
              WHERE major = 'CS')
```

Nested Query Example (Tricky)

- Find names of students applying to CS (using both join and nested queries).

MISS something, listen record again!

College(cName, state, enrollment)
Student(sID, sName, GPA, sizeHS)
Apply(sID, cName, major, decision)

```
SELECT DISTINCT sName  
FROM Student S, Apply A  
WHERE S.sID = A.sID AND A.major = 'CS'
```

if with DISTINCT will eliminate students with
same name but different ID
without DISTINCT, will have duplicates

```
SELECT sName  
FROM Student  
WHERE sID IN (SELECT sID  
              FROM Apply  
              WHERE major = 'CS')
```

Why are Duplicates Important?

- Find GPA of CS applicants (using both join and nested queries)

College(cName, state, enrollment)
Student(sID, sName, GPA, sizeHS)
Apply(sID, cName, major, decision)

```
SELECT GPA  
FROM Student  
WHERE sID in (SELECT  
sID FROM Apply WHERE  
major='CS')
```

SELECT GPA

Joins vs Sub-queries

- Many nested queries are equivalent to a simple query using JOIN operation. However, in many cases, the use of nested queries is necessary and cannot be replaced by a JOIN operation.
-
- **Advantages** of join queries
 - The join implementation can **display data from all tables** in the FROM clause whereas, the sub-query implementation can only **display data from the table(s)** in the outer query.
- **Advantages** of nested sub-queries
 - The ability to calculate an **aggregate value on the fly** and feed it back to the outer query for comparison is an advantage sub-queries have over joins.
- **Conclusion**
 - Use joins when you are displaying results from multiple tables.
 - Use sub-queries when you need to compare aggregates to other values.

Nested Queries in SQL

- Concept of nested sub-queries
- Correlated and non-correlated variants
- Sub-query operators
- Sub-queries vs. set operations and joins
- Division in SQL

Division in SQL

- Division in SQL is useful for answering queries include a “**for all**” or “**for every**” phrase, e.g., Find movie stars who were in all movies.
- Unfortunately, there is no direct way to express division in SQL. We can write this query, but to do so, we will have to express our query through double negation and existential quantifiers.

Examples of Division A/B

A

sno	pno
s1	p1
s1	p2
s1	p3
s1	p4
s2	p1
s2	p2
s3	p2
s4	p2
s4	p4

B1

pno
p2

B2

pno
p2
p4

B3

pno
p1
p2
p4

$A/B1$
the value of sno, with pno = p2

sno
s1
s2
s3
s4

$A/B2$

sno
s1
s4

$A/B3$

sno
s1

Division in SQL Using EXCEPT

- Find the IDs of movie stars who have played in all of the movies.

Movie(MovieID, Title, Year)
StarsIn(MovieID, StarID, Role)
MovieStar(StarID, Name, Gender)

SELECT StarID
FROM MovieStar MS
WHERE NOT EXISTS
All movies ((SELECT MovieID
FROM Movies)
Movies EXCEPT
Played by MS (SELECT MovieID
FROM StarsIn S
WHERE S.StarID=MS.StarID))

if the inner query
is empty means
that actor plays
in all the movies

BUT not supported
by SQL

Division in SQL Using NOT EXISTS

- Find the IDs of movie stars who have played in all of the movies.

```
Movie(MovieID, Title, Year)  
StarsIn(MovieID, StarID, Role)  
MovieStar(StarID, Name, Gender)
```

select Movie Star MS such that
there is no Movie M...
which is not played by MS

```
SELECT StarID  
FROM MovieStar MS  
WHERE NOT EXISTS  
(SELECT M.MovieID  
FROM Movie M  
WHERE NOT EXISTS  
(SELECT S.MovieID  
FROM StarsIn S  
WHERE S.MovieID=M.MovieID  
AND S.StarID=MS.StarID))
```

CREATE, ALTER and DROP TABLE statements

Basic SELECT Query

Set Operations

Aggregation, GROUP BY and HAVING

Nested Queries

Views

Null Values and Joins

INSERT, DELETE and UPDATE statements

Constraints and Triggers

Motivating Example for Use of Views

- Find those majors for which their average age is the minimum over all majors.

```
SELECT major, avg(age)
FROM student S
GROUP BY major
HAVING min(avg(age))
```

Wrong, cannot use
nested aggregation

- One solution would be to use subquery in the FROM Clause.

```
SELECT Temp.major, Temp.average
FROM(SELECT S.major, AVG(S.age) as average
      FROM Student S
      GROUP BY S.major) AS Temp
WHERE Temp.average in (SELECT MIN(Temp.average) FROM Temp)
```

Hideously ugly
Not supported
in all systems

- A Better alternative is to use views

What Are Views?

- A View is a single table that is derived from other tables, which could be base tables or previously defined views
- Views can be
 - **Virtual** tables - that do not physically exist on disk.
 - **Materialized** - by physically creating the view table.
These must be updated when the base tables are updated
- We can think of a virtual views as a way of specifying a table that we need to reference frequently, even though it does not physically exist.

Benefits of Using Views

- **Simplification:** View can hide the complexity of underlying tables to the end-users.
- **Security:** Views can hide columns containing sensitive data from certain groups of users.
- **Computed columns:** Views can create computed columns, which are computed on the fly.
- **Logical Data Independence:** Views provide support for logical data independence, that is users and user's programs that access the database are immune from changes in the logical structure of the database.

Defining and Using Views

```
CREATE VIEW <view name>  
  (<column name> {, <column name>} ) AS  
    <select statement> ;
```

- Example: Suppose we have the following tables:
 - Course(Course#,title,dept)
 - Enrolled(Course#,sid,mark)

```
CREATE VIEW CourseWithFails(dept, course#, mark) AS  
  SELECT C.dept, C.course#, mark  
  FROM Course C, Enrolled E  
  WHERE C.course# = E.course# AND mark<50
```

- This view gives the dept, course#, and marks for those courses where someone failed

Views and Security

- Views can be used to present necessary information (or a summary), while hiding details in underlying relation(s).

```
Course(Course#,title,dept)
Enrolled(Course#,sid,mark)
VIEW CourseWithFails(dept, course#, mark)
```

- Given CourseWithFails, but not Course or Enrolled, we can find the course in which some students failed, but we can't find the students who failed.

View Updates

- View updates must occur at the base tables.
 - Ambiguous
 - Difficult
- DBMS's restrict view updates only to some simple views on single tables (called updatable views)

Course(Course#,title,dept)
Enrolled(Course#,sid,mark)

VIEW CourseWithFails(dept, course#, mark)

Example: UQ has one table for students. Should the CS Department be able to update CS students info? Yes, Biology students? NO
Create a view for CS to only be able to update CS students.

Dropping Views

```
DROP VIEW [IF EXISTS]
view_name [, view_name] ...
[RESTRICT | CASCADE]
```

- Dropping a view does not affect any tuples of the underlying relation.
- DROP TABLE command has options to prevent a table from being dropped if views are defined on it:
 - RESTRICT : drops the table, unless there is a view on it
 - CASCADE: drops the table, and recursively drops any view referencing it

The Beauty of Views

- Find those majors for which their average age is the minimum over all majors.

```
SELECT Temp.major, Temp.average  
FROM(SELECT S.major, AVG(S.age) as average  
      FROM Student S  
     GROUP BY S.major) AS Temp  
 WHERE Temp.average in (SELECT MIN(Temp.average) FROM Temp)
```

Hideously ugly
Not supported
in all systems

```
Create View Temp(major, average) as  
    SELECT S.major, AVG(S.age) AS average  
    FROM Student S  
   GROUP BY S.major;
```

```
Select major, average  
From Temp  
WHERE average = (SELECT MIN(average) from Temp)
```

Clicker Question on Views

Consider the following table and SQL queries:

```
CREATE VIEW V AS  
    SELECT a+b AS d, c  
    FROM R;
```

```
SELECT d, SUM(c)  
FROM V  
GROUP BY d  
HAVING COUNT(*) <> 1;
```

d	a	b	c
2	1	1	3
3	1	2	3
3	2	1	4
5	2	3	5
6	2	4	1
5	3	2	4
6	3	3	6

Identify, from the list below, a tuple in the result of the query:

- A. (2,3)
- B. (3,12)
- C. (5,9)
- D. All are correct
- E. None are correct

result
d sum(c)
3 7
5 9
6 7

CREATE, ALTER and DROP TABLE statements

Basic SELECT Query

Set Operations

Aggregation, GROUP BY and HAVING

Nested Queries

Views

Null Values and Joins

INSERT, DELETE and UPDATE statements

Constraints and Triggers

NULL Values

- A value of NULL indicates that the value is unknown.
- The predicate **IS NULL** (**IS NOT NULL**) can be used to check for null values.
- Example: Find all student names whose age is not known.

```
SELECT name  
FROM Student  
WHERE age IS NULL
```

Operations on NULL Values

- NULL requires a 3-valued logic using the truth value *unknown*:
 - OR: $(\text{unknown or true}) = \text{true}$, $(\text{unknown or false}) = \text{unknown}$
 $(\text{unknown or unknown}) = \text{unknown}$
 - AND: $(\text{true and unknown}) = \text{unknown}$, $(\text{false and unknown}) = \text{false}$,
 $(\text{unknown and unknown}) = \text{unknown}$
 - NOT: $(\text{not unknown}) = \text{unknown}$
- Comparisons between two null values, or between a NULL and any other value, return unknown because the value of each NULL is unknown.
 - E.g. $5 < \text{null}$ or $\text{null} <> \text{null}$ or $\text{null} = \text{null}$
- All aggregate operations except **count(*)** ignore tuples with null values on the aggregated attributes.

```
select count(*)  
from class
```

```
select count(fid)  
from class
```

Clicker NULL Query

Determine the result of:

```
SELECT COUNT(*), COUNT(Runs)
FROM Scores
WHERE Team = 'Carp'
```

Which of the following is in the result?

- A. (1,0)
- B. (2,0)
- C. (1,NULL)
- D. All of the above
- E. None of the above

Scores:			
Team	Day	Opponent	Runs
Dragons	Sun	Swallows	4
Tigers	Sun	Bay Stars	9
Carp	Sun	NULL	NULL
Swallows	Sun	Dragons	7
Bay Stars	Sun	Tigers	2
Giants	Sun	NULL	NULL
Dragons	Mon	Carp	NULL
Tigers	Mon	NULL	NULL
Carp	Mon	Dragons	NULL
Swallows	Mon	Giants	0
Bay Stars	Mon	NULL	NULL
Giants	Mon	Swallows	5

Joins in SQL

- A Join used to **combine related tuples** from two relations into a single tuple in a new (result) relation.
- Join operation is needed for **organizing a search space** of data. This is needed when information is contained in more than one relation.
- Join relations are specified in the **FROM Clause**. When two relations are combined for a search, we need to know how the relations are combined.
- Based on **Cartesian Product** (denoted as X) There are three types of Join operations:
 - Theta-Join
 - Equi-Join
 - Natural Join

Cartesian Product in SQL

- Every row in R1 is matched with every row in R2 to form tuples in the result relation. The schema of the result relation contains all the columns from R1 and all the columns from R2

MovieStar

StarID	Name	Gender
1	Harrison Ford	Male
2	Vivian Leigh	Female
3	Judy Garland	Female

StarsIn

MovielD	StarID	Character
1	1	Han Solo
4	1	Indiana Jones
2	2	Scarlett O'Hara
3	3	Dorothy Gale

```
SELECT * FROM
MovieStar, StarsIn
```

For $|R1| = m$, $|R2| = n$, the $|R1 \times R2| = m * n$

MovieStar x StarsIn

StarID1	Name	Gender	MovielD	StarID2	Character
1	Harrison Ford	Male	1	1	Han Solo
2	Vivian Leigh	Female	1	1	Han Solo
3	Judy Garland	Female	1	1	Han Solo
1	Harrison Ford	Male	4	1	Indiana Jones
...

Clicker Question

- Consider the following SQL query:

```
SELECT DISTINCT s1.sname, s1.age  
FROM student s1, student s2  
WHERE s1.age > s2.age
```

- This query returns
 - A: The name and age of one of the oldest student(s).
 - B: The name and age of all of the oldest student(s).
 - C: The name and age of all of the youngest student(s).
 - D: The name and age of all students that are older than the youngest student(s).
 - E: None of the above.

Theta-Join in SQL

- Theta-join is the most general type of join, which allows several logical operators $\{=, \neq, <, \leq, >, \geq\}$.
- Example: For each student, list the students who are older than him/her (i.e., the first student).

```
SELECT A.EName, B.EName  
FROM Student A, Student B  
WHERE A.age < B.age
```

Theta-Join Example

MovieStar

StarID	Name	Gender
1	Harrison Ford	Male
2	Vivian Leigh	Female
3	Judy Garland	Female

StarsIn

MovielID	StarID	Character
1	1	Han Solo
4	1	Indiana Jones
2	2	Scarlett O'Hara
3	3	Dorothy Gale

```
SELECT *
FROM MovieStar M, StarsIN S
Where M.StarID < S.StarID
```

1	Name	Gender	MovielID	5	Character
1	Harrison Ford	Male	2	2	Scarlett O'Hara
1	Harrison Ford	Male	3	3	Dorothy Gale
2	Vivian Leigh	Female	3	3	Dorothy Gale

Equi-Join and Natural Join

- Equi-Join: A special case of Theta join where condition contains only *equalities*.
- The SQL NATURAL JOIN is a type of Equi-Join and is structured in such a way that, columns with same name of associate tables will appear once only.
- Natural Join : Guidelines
 - The associated tables have one or more pairs of identically named columns.
 - The columns must be the same data type.

Natural join of tables with no pairs of identically named columns will return the cross product of the two tables.

Natural Join Examples

MovieStar

StarID	Name	Gender
1	Harrison Ford	Male
2	Vivian Leigh	Female
3	Judy Garland	Female

StarsIn

MovieID	StarID	Character
1	1	Han Solo
4	1	Indiana Jones
2	2	Scarlett O'Hara
3	3	Dorothy Gale

Select *
From MovieStar natural join StarsIN

StarID	Name	Gender	MovieID	Character
1	Harrison Ford	Male	1	Han Solo
1	Harrison Ford	Male	4	Indiana Jones
3	Judy Garland	Female	3	Dorothy Gale
2	Vivian Leigh	Female	2	Scarlett O'Hara

Inner and Outer Joins

- **Inner Join**

This is the default join, in which a tuple is included in the result relation, only if matching tuples exist in both relations.

- **Outer Join**

Outer joins can include the tuples that do not satisfy the join condition, i.e, a matching tuple does not exist, which is indicated by a NULL value

- Full Outer Join: Includes all rows from both tables.
- Left Outer Join: Includes all rows from first table.
- Right Outer Join: Includes all rows from second table .

```
SELECT <attribute list>
```

```
FROM table_reference {LEFT | RIGHT} [OUTER] JOIN table_reference ON  
< search_condition>
```

Inner and Outer Joins Examples

A	B
1	2
3	3

B	C
2	4
4	6

Natural
Inner Join

A	B	C
1	2	4

Natural
Left outer Join

A	B	C
1	2	4
3	3	Null

Natural
Right outer Join

A	B	C
1	2	4
Null	4	6

Natural
outer Join

A	B	C
1	2	4
3	3	Null
Null	4	6

Outer join (without the Natural) will use the key word on for specifying the condition of the join.

Outer join not implemented in MySQL
Outer join is implemented in Oracle

Clicker Outer Join Question

- Given the following relations Compute:

```
SELECT R.A, R.B, S.B, S.C, S.D  
FROM R FULL OUTER JOIN S  
ON (R.A > S.B AND R.B = S.C)
```

- Which of the following tuples of R or S is dangling (and therefore needs to be padded with NULLs in the outer join)?

- A. (1,2) of R
- B. (3,4) of R
- C. (2,4,6) of S
- D. All of the above
- E. None of the above

R(A,B)		S(B,C,D)		
A	B	B	C	D
1	2	2	4	6
3	4	4	6	8
5	6	4	7	9

CREATE, ALTER and DROP TABLE statements

Basic SELECT Query

Set Operations

Aggregation, GROUP BY and HAVING

Nested Queries

Views

Null Values and Joins

INSERT, DELETE and UPDATE statements

Constraints and Triggers

INSERT Statement

- INSERT statement is used to add tuples to an existing relation
- Single Tuple INSERT
 - Specify the relation name and a list of values for the tuple
 - Values are listed in the **same order** as the attributes were specified in the CREATE TABLE command
 - User may specify **explicit attribute names** that correspond to the values provided in the insert statement. The attributes not included cannot have the NOT NULL constraint
- Multiple Tuple INSERT
 - By separating each tuple's list of values with commas
 - By loading the result of a query

Single Tuple INSERT Example

INSERT INTO <table name>

[(<column name> {, <column name>})]

(VALUES (<constant value>, {,<constant value>})
| <select statement>);

- Can insert a single tuple using:

INSERT INTO Student

VALUES (53688, ‘Smith’, 3.2, 200)

- or

INSERT INTO Student (sID, sName, GPA, sizeHS)

VALUES (53688, ‘Smith’, 3.2, 200)

- Add a tuple to student with null address and phone:

INSERT INTO Student (sID, sName, GPA, sizeHS)

VALUES (53688, ‘Smith’, 3.2, NULL)

Multiple Tuple INSERT Example

INSERT INTO <table name>

[(<column name> {, <column name>})]

(VALUES (<constant value>, {,<constant value>})
| <select statement>);

- Can add values selected from another table
- Make student 123 apply into all “BIO” related majors at Stanford.

```
INSERT INTO apply
SELECT 123, "Stanford", major, NULL
FROM apply
WHERE major LIKE "%bio%"
```

The select-from-where statement is fully evaluated before any of its results are inserted or deleted.

DELETE Statement

- DELETE statement is used to remove existing tuples from a relation.
- A single DELETE statement may delete zero, one, several or all tuples from a table.
- Tuples are explicitly deleted from a single table.
- Deletion may **propagate to other tables** if referential triggered actions are specified in the referential integrity constraints of the CREATE (ALTER) TABLE statement

```
DELETE FROM <table name>  
[WHERE <select condition>];
```

DELETE Statement Example

- Delete all “BIO” related applications of student 123 for Stanford.

```
DELETE FROM apply  
WHERE sID = 123 AND  
cName LIKE "Stanford" AND major LIKE "%BIO%"
```

- Note that only whole tuples are deleted.
- Can delete all tuples satisfying some condition

UPDATE Statement

UPDATE statement is used to modify attribute values of one or more selected tuples in a relation.

- Tuples are selected for update from a single table.
- However, updating a primary key value may **propagate to other tables** if referential triggered actions are specified in the referential integrity constraints of the CREATE (ALTER) TABLE statement.

```
UPDATE <table name>
    SET <column name> = <value expression>
        {, <column name> = <value expression>}
    [WHERE <select condition>];
```

UPDATE Statement Example

- Increase the age of all students by 2 (should not be more than 100)
- Need to write two updates:

```
UPDATE Student  
SET      age = 100  
WHERE    age >= 98
```

```
UPDATE Student  
SET age = age + 2  
WHERE age < 98
```

- Is the order important?

CREATE, ALTER and DROP TABLE statements

Basic SELECT Query

Set Operations

Aggregation, GROUP BY and HAVING

Nested Queries

Views

Null Values and Joins

INSERT, DELETE and UPDATE statements

Constraints and Triggers

Semantic Constraints

- Keys, entity constraints and referential integrity are structural constraints that are managed by the DBMS.
- Semantic constraints can be specified using CHECK and ASSERTION statements.
- The constraint is satisfied by a database state if no combination of tuples in the database state violates the constraint.

Semantic Constraints: Check

- Semantic constraints over a single table are specified using tCheck conditional-expressions

```
CREATE TABLE Student
( snum INTEGER,
  sname CHAR(32),
  major CHAR(32),
  standing CHAR(2)
  age REAL,
  PRIMARY KEY (snum),
  CHECK ( age >= 10
    AND age < 100 );
```

Check constraints are checked when tuples are inserted or modified

Constraints Over Multiple Relations

- Constraints that cannot be defined in one table are defined as ASSERTIONS which are not associated with any table.

Example: *Every MovieStar needs to star in at least one Movie.*

Movie(MovieID, Title, Year)
StarsIn(MovieID, StarID, Role)
MovieStar(StarID, Name, Gender)

```
CREATE ASSERTION totalEmployment
CHECK
( NOT EXISTS (
    SELECT StarID
    FROM MovieStar
    WHERE StarID not in ( SELECT StarID
                           FROM StarsIn)));
```

Triggers

- An active database is a database that includes an event-driven architecture. Triggers are a procedure that start automatically if specified changes occur to the DBMS.
- A trigger has three parts:
 1. Event (activates the trigger)
 2. Condition (tests whether the trigger should run)
 3. Action (procedure executed when trigger runs)
- Database vendors did not wait for trigger standards! So trigger format depends on the DBMS

NOTE: triggers may cause cascading effects.

Triggers: Example (SQL:1999)

```
CREATE TRIGGER youngStudentUpdate  
    AFTER INSERT ON Student
```

event

```
REFERENCING NEW TABLE NewStudent  
FOR EACH STATEMENT
```

newly inserted
tuples

```
INSERT INTO
```

apply once per
statement

action

```
YoungStudent(snum, sname, major, standing, age)
```

```
SELECT snum, sname, major, standing, age
```

```
FROM NewStudent N
```

```
WHERE N.age <= 18;
```

Can be either before or after

Learning Objectives Revisited

Description	Tag
Create basic SQL queries using: SELECT, FROM, WHERE statements.	
Create SQL queries containing the DISTINCT statement.	
Create SQL queries using set operators.	
Create SQL queries using aggregate operators.	SQL-basic
Create SQL queries containing GROUP BY statements.	
Create SQL queries containing HAVING statements.	
Given a SQL query and table schemas and instances, compute the query result.	
Create nested SQL queries.	
Create SQL Queries that use the division operator.	
Explain the purpose of NULL values, and justify their use. Also describe the difficulties added by having NULLs .	
Create SQL queries that use joins.	
Create SQL queries that use the CHECK statement.	
Create SQL queries that use ASSERTIONS .	
Create, use and modify VIEWS in SQL.	SQL-advance
Modify data stored in a database using the INSERT, DELETE, and UPDATE statements.	
Identify the pros and cons of using general table constraints (e.g., ASSERTION, CHECK) and triggers in databases.	
Show that there are alternative ways of coding SQL queries to yield the same result.	
Determine whether or not two SQL queries are equivalent.	
Create SQL queries for creating tables.	
Create SQL queries for altering tables.	
Create SQL queries to enforce referential integrities.	
Create SQL queries for dropping tables.	SQL-DDL